

JUAN J. GOMEZ

Tampa, FL · (813) 304-9152 · jjgomezswe@gmail.com · linkedin.com/in/jjgomezswe · github.com/jjosegomez · Permanent US Resident (no sponsorship) · Bilingual EN/ES

SUMMARY

AI-native software engineer who **architected and shipped a Fortune-500 finance org's first production AI agent — solo** — serving 6 global analysts at **96% accuracy**, while concurrently delivering three other production apps. Specializes in **agentic AI**: LLM agent design, evaluation frameworks, RAG, and AI-assisted software delivery. Daily driver of Claude, Cursor, and Kiro for agent-assisted development; builds the eval suites and quality gates that validate agent-generated output before production.

TECHNICAL SKILLS

Languages: SQL, Python, TypeScript, JavaScript, Java, C#, C++

Agentic AI / LLM: Snowflake Cortex (Agent/Analyst), AWS Bedrock, Anthropic Claude, OpenAI, LLM agent design & orchestration, RAG, prompt engineering & context design, AI evaluation frameworks, DeepEval

AI-Assisted Dev: Kiro, Claude / Claude Code, GitHub Copilot, Cursor — agent-assisted delivery + review of AI-generated code/tests for quality & alignment

Frameworks & Cloud: Next.js, FastAPI, Streamlit, React, Node.js, Supabase/PostgreSQL, Snowflake, AWS (Bedrock, EC2, S3), Azure (App Service, DevOps/CI-CD), Docker, Okta

Quality & Delivery: automated testing (Playwright, Vitest), AI eval suites, code review, release validation, rollback/credential-rotation runbooks, row-level access, Agile/Scrum, Git

EXPERIENCE

Jabil — Financial Systems Developer II

Tampa, FL · Apr 2025 – Present

Flagship: Jabil Finance's first production AI agent

- ▶ **Sole architect & developer of Jabil Finance's first production AI agent** (Snowflake Cortex) — 9 views, 8 semantic models, orchestration instructions, and an evaluation framework — serving **6 global analysts at 96% accuracy** against SAP/Cognos source reports. First-mover at Jabil.
- ▶ Built an **AI evaluation suite (53+ test cases)** validating agent accuracy vs. source-of-truth financial reports — the quality gate gating every release.
- ▶ Authored **A.G.E.N.T.**, a repeatable methodology/playbook for deploying LLM agents — now positioned as internal consultant as other teams adopt agentic systems.
- ▶ Shipped **three additional production/MVP apps** concurrently: a Pillar II tax app (v1.1 to prod, **zero cutover incidents**, 10,835-row audit migration); a File Manager with **AI metadata extraction (Claude Haiku via AWS Bedrock)**; and a Next.js/FastAPI finance platform prototype.
- ▶ Own **end-to-end production delivery**: release checklists, credential rotation, rollback criteria, row-level access policies — full-stack (Python/Streamlit, Next.js/FastAPI), Okta auth, Docker, Azure CI/CD.
- ▶ Work **directly with cross-functional finance stakeholders** (no intermediary) on requirements, validation, and data-discrepancy resolution.

Medtronic — Software Engineer I

Minneapolis, MN · Mar 2024 – Apr 2025

- ▶ Developed full-stack apps and manufacturing-equipment integrations in **Python** and Ignition, automating data collection via the **SECS/GEM** protocol.
- ▶ Built a custom **OPC UA device module in Java** (Ignition SDK) for seamless equipment-to-software communication across all connection types.
- ▶ Architected a **Microsoft SQL Server** redesign (normalization) for query performance; led **Validation Plans** to FDA/Medtronic quality standards; managed DevOps via Git/GitLab.

Eversana Intouch — Software Development Intern

Kansas City, KS · Jun – Aug 2023

- Redesigned live client sites (HTML/CSS/JS) for responsiveness & accessibility; built a strategic **React + Node.js** app that lifted traffic and retention across three concurrent projects.

FLAGSHIP FOUNDER PROJECT — TRINITY (VIGO LEGAL)

Founding Engineer · 2025 – Present · production multi-agent AI legal platform, live pilot firms

- Architected a **production multi-agent platform**: intake classification → RAG over Colombian jurisprudence (tri-source) → grounded legal drafting under a strict **propose-then-approve** human-in-the-loop control.
- Built an **LLM evaluation harness** (DeepEval, 13 cases × 6 metrics) gating every retrieval/prompt change; practice **subagent-driven development** with adversarial **review of agent-generated code** (caught/fixed 7 pre-merge defects); hardened untrusted-doc ingestion against prompt injection.

EDUCATION

- **Georgia Institute of Technology** — M.S. Computer Science (OMSCS), Machine Learning specialization · *Starting Fall 2026*
- **University of South Florida** — B.S. Computer Science · 2020–2023 · GPA 3.49

INVOLVEMENT

- **Google Developer Student Club, USF** — Tech Lead (2023–2024): 10+ workshops (500+ students), mentored 10+ students, ran university hackathons. **SHPE** — Mentor.